

Am I the Ahole?: Do LLMs reproduce community-endorsed moral advice?**

Theoretical Framework

A central question underlying this paper is what it means to reason morally in the first place. The issue is especially important when evaluating large language models (LLMs), whose outputs often appear thoughtful, balanced, and ethically sensitive. The production of morally acceptable judgments, however, is not necessarily the same thing as moral reasoning. A system may consistently generate responses that resemble moral deliberation without possessing the underlying capacities that philosophers and psychologists typically associate with moral agency.

Several frameworks from moral philosophy suggest that moral reasoning involves more than isolated judgments. Peter Railton argues that moral evaluation requires movement beyond immediate intuitions toward principles that can be generalized across cases. Likewise, Bernard Williams distinguishes between merely following external rules and acting from reasons that are meaningfully connected to an agent's motivations and commitments. In these views, moral reasoning is not simply arriving at the "right answer." It involves integrating particular judgments into a broader structure of values, reasons, and justificatory practices. Research in moral psychology similarly emphasizes the importance of consistency and reflection. Michael Tomasello argues that human morality emerges through social cooperation and shared normative expectations, allowing individuals to evaluate both their own actions and those of others. Julia Haas's framework for evaluating moral competence likewise highlights the importance of applying moral concepts across situations rather than treating each judgment as an isolated event. Even when people disagree morally, they typically understand themselves as participating in a larger system of norms that extends across time and contexts.

This continuity matters because moral reasoning is not merely a one-time verdict. It is part of an ongoing process of reflection, self-correction, and negotiation between competing values. Humans revise beliefs, reconcile contradictions, and try to bring their intuitions into coherence over time. By contrast, large language models have no enduring moral self, no stable normative commitments, and no intrinsic investment in the ethical principles they articulate. Their outputs are generated through probabilistic optimization over linguistic patterns rather than participation in moral life itself. A model may explain utilitarianism persuasively in one context and then

adopt a different moral framework in another, depending on the wording of the prompt, the emotional tone of the user, or the kind of answer the interaction seems to call for.

This raises the possibility that discussions of “AI morality” involve a category mistake. We may be mistaking the successful simulation of moral discourse for moral cognition itself. A language model can imitate the vocabulary of ethical reflection without possessing the features that make moral reasoning meaningful in humans: agency, accountability, continuity, or genuine normativity. In this sense, the appearance of ethical sophistication may reflect alignment training and contextual adaptation more than principled moral deliberation.

Accordingly, this paper evaluates LLM moral judgment by comparing model verdicts to human-consensus judgments in interpersonal moral dilemmas. Using Reddit AITA verdicts as a community-based baseline, we examine whether ChatGPT and Claude reproduce human patterns of moral evaluation across cases involving blame, ambiguity, mutual fault, and insufficient information. This approach does not treat Reddit consensus as objective moral truth. Rather, it uses human agreement as a comparison point for analyzing how model judgments align with, diverge from, or flatten common forms of interpersonal moral reasoning. If models produce plausible moral explanations while showing uneven agreement across verdict categories, avoiding uncertainty, or diverging sharply from one another, then their outputs may reveal less about stable moral reasoning and more about the dynamics of language modeling, alignment, and context-sensitive response generation.

Benchmarking Moral Reasoning

Evaluating moral judgment presents a challenge that differs from traditional benchmarking tasks in machine learning. Metrics such as accuracy, coherence, and helpfulness are useful for measuring general language performance, but they provide only a partial picture of moral evaluation. A model may generate responses that appear thoughtful, balanced, and ethically sophisticated while still diverging from the ways humans assign blame, distribute responsibility, or navigate uncertainty. Assessing moral reasoning therefore requires examining not only the quality of a model’s explanations, but also the structure of its judgments relative to human evaluative practices.

To do so, this paper uses Reddit’s *Am I the Asshole?* (AITA) verdicts as a human-consensus baseline for evaluating LLM moral judgment. The goal is not to treat Reddit users as objective moral authorities, but rather, to use their aggregated verdicts as a comparison point for common interpersonal moral evaluation. Users present disputes involving family relationships, friendships, romantic partnerships, workplace conflicts, questions of fairness, and competing obligations. Other users then evaluate these situations using a shared verdict system consisting of **YTA** (You’re the Asshole), **NTA** (Not the Asshole), **ESH** (Everyone Sucks Here), **NAH** (No Assholes Here), and **INFO** (Not Enough Information).

This kind of benchmark allows us to ask a more specific question: when LLMs are asked to evaluate interpersonal moral dilemmas, do their verdicts align with human consensus, and where do they systematically diverge? Overall agreement rates can show whether models broadly reproduce human moral judgments, while category-level agreement can reveal whether models handle some kinds of cases better than others. For example, a model may perform well on clear blame cases but struggle with cases involving mutual responsibility or insufficient information. Directional disagreement analysis can also reveal whether a model tends to be more lenient or more severe than human raters.

In this sense, benchmarking LLM moral judgment is a way of testing whether fluent moral language corresponds to stable and human-recognizable patterns of moral evaluation. If a model gives plausible explanations but avoids uncertainty, overassigns blame, underrecognizes mutual fault, or diverges sharply from both human consensus and other models, then its outputs may reflect conversational optimization more than genuine moral reasoning. Agreement with human consensus does not prove that a model possesses moral understanding, but disagreement patterns can help reveal the limits of treating moral fluency as moral cognition.

Experimental Design

To evaluate how large language models respond to interpersonal moral dilemmas, we compare model-generated verdicts to human-consensus judgments from the AITA subreddit. Rather than using abstract philosophical thought experiments alone, this design focuses on everyday moral conflicts involving blame, responsibility, ambiguity, and interpersonal harm. AITA posts are useful for this purpose because each submission presents a real or realistic moral dilemma, and

the comment section produces a collective judgment through verdict labels such as YTA, NTA, ESH, NAH, and INFO.

Our dataset consists of 300 AITA submissions, balanced across the five verdict categories. Each category contains 60 posts. This balanced sampling strategy prevents the analysis from being dominated by whichever verdicts are most common on the subreddit and allows category-specific performance to be examined directly. Each category represents different forms of moral evaluation: YTA/NTA involve relatively clear judgments of blame or justification, ESH captures mutual wrongdoing, NAH reflects conflicts in which neither party is considered morally blameworthy, and INFO functions as a judgement about the inadequacy of evidence.

For each submission, the title and body text were presented to two LLMs, ChatGPT and Claude, and each model was asked to assign one of the five AITA verdict categories alongside an explanation for its decision. These responses were then compared to the human-consensus verdict associated with each post. This design allows us to measure not only overall agreement with human judgment but also variation across verdict categories, disagreement patterns, and differences between models.

The goal of the experiment is not to determine the objectively correct moral answer to each dilemma. Reddit consensus is not treated as moral truth, but rather as a community-based baseline for comparison. By examining where models align with or diverge from this human baseline, we can better understand the kinds of moral judgments they reproduce, the kinds they systematically reinterpret, and the conditions under which model-generated moral reasoning differs from human moral reasoning.

Expected Findings

We hypothesize that model performance will vary across verdict categories rather than be uniformly aligned with human consensus. In particular, we anticipate stronger agreement on relatively clear cases of blame or justification, and more divergence on cases involving ambiguity, shared responsibility, or informational uncertainty (ESH, NAH, INFO).

We expect the greatest divergence to appear in INFO cases. Human evaluators label a post as INFO when they are identifying the situation as under-specified and requiring more context

before moral responsibility can be assigned. Since LLMs are optimized to produce determinate and helpful answers, we predict that both ChatGPT and Claude will underutilize the INFO category and instead assign more definitive verdicts. Such a pattern would suggest that the models are more comfortable resolving uncertainty than preserving it.

We also expect model judgments to differ from each other, rather than reflecting one unified pattern of “AI moral reasoning.” Because different models are shaped by different training data, alignment procedures, and response norms, ChatGPT and Claude may show different tendencies in how they assign blame. One model may be more likely to frame conflicts as mutual fault, while another may be more likely to soften blame or avoid harsh conclusions. These differences would suggest that LLM moral judgment is not monolithic, but model-specific and shaped by optimization processes.

Data Preprocessing

The Reddit Dataset for AmItheAsshole (AITA) Subreddit contains both submissions and their corresponding comment threads. Each submission presents a personal moral dilemma, and community members respond by labeling the original poster (OP) with one of five verdicts:

- YTA (You're the A-hole) – The OP is primarily at fault in the situation
- NTA (Not the A-hole) – The OP is not at fault; their behavior was reasonable
- ESH (Everyone Sucks Here) – Multiple parties, including the OP, behaved poorly
- NAH (No A-holes Here) – No one is truly at fault; the situation is a genuine conflict without clear wrongdoing
- INFO (Not Enough Information) – Commenters need more context before making a judgment

For our dataset, we selected 300 submissions and aggregated all corresponding top-level comments to extract verdict labels. To assign a single ground-truth label to each submission, we computed the mode of all verdicts across the comment thread. Therefore, the most frequently occurring category was treated as the consensus community opinion. This resulted in a labeled dataset of 300 submissions with human-consensus ground-truth labels for downstream evaluation.

Evaluation Metrics

1. Verdict Distribution

The first metric examines the overall distribution of verdicts assigned by each source. By comparing the proportion of YTA, NTA, ESH, NAH, and INFO verdicts across human raters, ChatGPT, and Claude, we can identify systematic tendencies in how each source approaches moral judgment.

2. Overall Agreement Rate.

The primary agreement metric measures the proportion of cases in which each model's verdict matches the human consensus verdict. This is computed separately for ChatGPT and Claude, as well as jointly to identify cases where both models agree with humans simultaneously and cases where neither does. This metric provides a baseline measure of alignment but is treated as a starting point rather than a conclusion, since agreement alone does not capture the nature or direction of disagreement.

3. Agreement Rate by Verdict Category

Overall agreement rates can obscure meaningful variation across verdict types. We therefore disaggregate agreement by human verdict category, measuring how often each model correctly identifies YTA, NTA, ESH, NAH, and INFO cases respectively. This allows us to assess whether models perform consistently across the moral spectrum or whether alignment is stronger in some categories than others..

4. Directional Disagreement Analysis

When a model's verdict diverges from the human consensus, the direction of that divergence carries interpretive significance. We distinguish between two primary directional biases: leniency bias, defined as cases where the human verdict is YTA but the model assigns NTA, and severity bias, defined as the reverse. Cross-tabulation of human verdicts against model verdicts for all disagreement cases allows us to map the full landscape of how models misclassify moral situations and whether that misclassification follows a consistent pattern.

5. Inter-Model Agreement

Beyond comparing each model to the human baseline, we also measure the degree to which ChatGPT and Claude agree with each other independently of human consensus. Cases where the two models diverge from one another — even when one is aligned with the human verdict — reveal differences in the underlying moral reasoning of each system and suggest that LLM moral judgment is not monolithic.

6. How Models Classified INFO Cases

Because INFO represents cases where human commenters judged that there was not enough information to assign responsibility, we separately examine how each model reclassified these cases when it did not choose INFO. This metric focuses only on posts where the human-consensus verdict was INFO and compares the verdicts assigned by ChatGPT and Claude. By doing so, we can identify whether models preserve uncertainty or instead force under-specified dilemmas into more determinate categories such as YTA, NTA, ESH, or NAH.

7. Average Moral Theme Rates: OpenAI vs. Claude

In addition to comparing verdict labels, we also examine the moral language used by ChatGPT and Claude in their explanations. We calculate average moral theme rates by measuring the frequency of keywords related to themes such as care and harm, fairness and reciprocity, loyalty and family, authority and rules, autonomy and consent, proportionality, and recognition and dignity. This allows us to compare not only what verdict each model assigns, but also what kinds of moral concerns each model emphasizes while justifying its decision.

Results




Figure 1: Verdict Distribution Between Human vs. ChatGPT vs. Claude

Verdict Distribution (%)			
	Human	ChatGPT	Claude
YTA	20.0	21.0	26.7
NTA	20.0	24.0	27.0
ESH	20.0	28.7	22.3
NAH	20.0	26.0	23.3
INFO	20.0	0.3	0.7

Table 1: Percentage Verdict Breakdown Between Human vs. ChatGPT vs. Claude

Agreement with Human Verdict (n = 300)	
ChatGPT agrees with human	200 (66.7%)
Claude agrees with human	136 (45.3%)
Both agree with human	127 (42.3%)
Neither agree with human	91 (30.3%)

Table 2: LLM agreement with human verdict

Agreement Rate by Human Verdict Category			
Human Verdict	ChatGPT Agrees (%)	Claude Agrees (%)	Both Agree (%)
YTA	76.7	68.3	60.0
NTA	88.3	71.7	71.7
ESH	100.0	33.3	33.3
NAH	66.7	51.7	46.7
INFO	1.7	1.7	0.0

Table 3: Agreement Rate By Human Verdict Category




Figure 2: Model Agreement with Human Verdict By Category




Figure 3: What ChatGPT Says Instead




Figure 4: What Claude Says Instead

OpenAI Leniency & Severity Bias	
Leniency bias (human=YTA/ESH, model=NTA/NAH)	7 / 120 (5.8%)
Severity bias (human=NTA/NAH, model=YTA/ESH)	17 / 120 (14.2%)

Table 4: OpenAI/ChatGPT Leniency and Severity Bias

Leniency breakdown:

- Human=YTA, OpenAI=NAH: 7

Severity breakdown:

- Human=NTA, OpenAI=ESH: 5
- Human=NAH, OpenAI=YTA: 5
- Human=NAH, OpenAI=ESH: 7

Claude Leniency & Severity Bias	
Leniency bias (human=YTA/ESH, model=NTA/NAH)	33 / 120 (27.5%)
Severity bias (human=NTA/NAH, model=YTA/ESH)	30 / 120 (25.0%)

Leniency breakdown:

- Human=YTA, Claude=NTA: 1
- Human=YTA, Claude=NAH: 7
- Human=ESH, Claude=NTA: 13
- Human=ESH, Claude=NAH: 12

Severity breakdown:

- Human=NTA, Claude=ESH: 11

- Human=NAH, Claude=YTA: 8
- Human=NAH, Claude=ESH: 11




Figure 5: ChatGPT vs. Claude Verdict Agreement




Figure 6: How Models Classified INFO Cases




Figure 7: Average Moral Theme Rates: OpenAI vs. Claude

Discussion

Our results show that the two models reproduce human consensus to very different degrees. As Table 2 reports, ChatGPT matched the human verdict in 200 of 300 cases (66.7%), while Claude matched in only 136 (45.3%). Read on its own, this gap looks like one model simply outperforming the other, yet, the category-level results in Table 3 and the disagreement heatmaps in Figures 3 and 4 tell a more interesting story of how the two models are not better and worse versions of the same judgment process, but different judgment processes with different default tendencies. The verdict distributions in Figure 1 and Table 1 represent how ChatGPT's verdicts cluster in the middle of the moral spectrum, with ESH (28.7%) and NAH (26.0%) as its two most common outputs, both above the human rate of 20%. Claude's verdicts cluster at the edges, with YTA (26.7%) and NTA (27.0%) as its two most common outputs. In other words, ChatGPT tends to reach for "everyone is at fault" or "no one is at fault," while Claude tends to polarize, deciding that one party is clearly responsible or clearly not. The disagreement heatmaps make this concrete. In Figure 4, when the human consensus was ESH, Claude almost never agreed; it

reassigned those cases to YTA 15 times, NTA 13 times, and NAH 12 times, leaving only 20 of 60 as ESH. Claude is actively breaking shared-fault situations down into a single responsible party or no responsible party. Figure 5 shows the same thing model to model: of the 86 cases where ChatGPT said ESH, Claude agreed on only 35 and reassigned the other 51 to single-party or no-fault verdicts.

Further, in Table 3, Claude agreed with the human ESH verdict in only 33.3% of ESH cases, while ChatGPT agreed in 100%. Looping back, Figure 1 shows ChatGPT assigns ESH far more often than humans do, roughly 86 times across the dataset when only 60 cases are human-defined ESH. It reaches a perfect agreement rate within the ESH category largely by applying ESH broadly, which means catching all 60 true cases at the cost of about 26 false ones. Its ESH precision is therefore only around 70%, even though its recall is perfect. The 100% in Table 3 reflects how readily ChatGPT reaches for "everyone sucks here," not necessarily a fine-grained sensitivity to shared fault, and the category agreement rate captures how often a model finds the true cases but not how often it is wrong when it makes that call.

The two models also handle uncertainty differently. Figure 1 shows that ChatGPT and Claude each use INFO under 1% of the time against a human rate of 20%, and Table 3 shows both agree with human INFO verdicts only 1.7% of the time. But Figures 3, 4, and 6 show they dispose of those cases differently. When humans said INFO, ChatGPT pushed 29 of 60 cases into NAH, defaulting uncertainty toward "no one is at fault," consistent with its broader tendency to mutualize fault. Claude scattered the same cases almost evenly across YTA (16), NTA (15), ESH (14), and NAH (14), forcing a determinate verdict in whatever direction the case leaned. Figure 6 makes this difference especially clear because ChatGPT has one very visible spike at NAH, while Claude spreads its reclassifications more evenly across the other categories.

The directional results in Table 4 follow from these tendencies rather than standing apart from them. ChatGPT shows low leniency bias (5.8%) and moderate severity bias (14.2%), which fits a model that mostly stays near human consensus and errs toward the shared-fault middle. Claude shows elevated rates of both leniency (27.5%) and severity (25.0%) at once, which is exactly what we would expect from a model pulling cases out of the ambiguous middle in both directions. Its leniency comes almost entirely from softening ESH into NTA or NAH (25 of its

33 lenient cases), and its severity comes from hardening NTA and NAH into ESH or YTA. The two biases are not in tension with each other because they both come from the same broader tendency to avoid leaving a situation in a mutually blameworthy or under-specified middle.

Figure 7 adds another layer to this interpretation by showing that the models also differ in the kinds of moral language they use, not just in the verdict labels they assign. This is because across the moral theme categories, both models rely most heavily on loyalty and family language, but ChatGPT does so at a higher rate than Claude. ChatGPT also uses more language related to care and harm, authority and rules, autonomy and consent, and recognition and dignity, while Claude is slightly higher only in fairness and reciprocity and proportionality. This suggests that ChatGPT's pull toward ESH and NAH is not only a verdict-level pattern, but also appears in the language it uses to frame moral problems. It seems more likely to talk through relational harm, responsibility, boundaries, and dignity, while Claude's verdict pattern suggests a stronger tendency to decide where blame should land. In that sense, the models are not only making different decisions; they are emphasizing different moral vocabularies while making those decisions.

These patterns matter because they show that the models are not just randomly disagreeing with human consensus and have structure on their own. From the results, Claude does not avoid ESH at random as it repeatedly breaks shared-fault situations into either a single responsible party or no responsible party at all. ChatGPT, by contrast, is much more willing to place cases in the middle by assigning ESH or NAH. Figure 7 reinforces this broader point because the models also differ in the moral themes they emphasize, meaning their divergence is not only numerical but also rhetorical and interpretive. This complicates the argument of the paper in an important way, because if a human showed this kind of stable pattern across cases, we might be tempted to read it as a real moral tendency or evaluative commitment. In that sense, the models are not just producing disconnected moral outputs but recognizable judgment profiles.

However, a recognizable pattern is still not the same thing as moral reasoning because models can apply the same tendency across cases without being accountable to that tendency as a principle. It does not remember its past judgments in the way a person does, does not experience guilt or contradiction, and does not feel any need to explain why it moved away from a previous

commitment. This is why consistency alone is not enough. The regularity we observe here happens across cases, but it is not part of an ongoing process of reflection, revision, or self-correction. This distinction is where the difference between aligned AI and ethical AI becomes important. ChatGPT's tendency toward mutual fault and Claude's tendency toward decisive single-party verdicts look like different moral temperaments, but they are better understood as products of training, alignment, and response style rather than moral theories the models actually hold. The danger is that the surrounding explanations sound so fluent and thoughtful that it becomes easy to treat these patterns as evidence of moral character, but the models are not endorsing principles in the human sense as they are generating judgments that are patterned, model-specific, and often persuasive, while still lacking the agency, accountability, and continuity that make moral reasoning meaningful.

So the results do not show that LLM moral judgments are meaningless, but they also do not show that the models are genuinely reasoning morally. What they show is that ChatGPT and Claude produce moral judgments with identifiable patterns, and those patterns reveal the effects of model-specific optimization more than stable normative commitment. This is why moral fluency remains risky to overinterpret. The models can sound like they are reasoning through a dilemma, and sometimes they even agree with human consensus, but their outputs still reflect the simulation of moral judgment rather than moral cognition itself.

Limitations

Several features of our design limit how far these results can be pushed. First, we treated the mode of community verdicts as the ground-truth label, which discards information about how divided the community was. A post where 95% of commenters said YTA and a post where 51% did are treated identically, even though a model that disagrees with the second is disagreeing with a community that was itself split. Our agreement rates therefore penalize models on genuinely contested cases as heavily as on clear ones.

Second, our dataset is balanced at 60 posts per category, which does not reflect the natural distribution of AITA verdicts, where NTA dominates and INFO is rare. Balancing lets us compare performance across categories, but it means our overall agreement rates are not estimates of how often these models would agree with humans on real subreddit traffic. The

INFO category in particular is upsampled far above its natural frequency, so a category the models almost never use accounts for a fifth of the cases behind the headline agreement numbers. Reported with INFO excluded, both models' agreement rates rise.

Third, the INFO section is a limitation because for a human commenter, withholding judgment opens a dialogue that the original poster can answer, whereas a model asked for a single verdict has no such follow-up available and little reason to decline. The models' near-total avoidance of INFO is consistent with a bias toward decisive answers, but it is also partly an artifact of asking for one verdict in one turn but this in turn gave a limitation that there was a lack of usage of the INFO category by the models to further explore and reason through. Potentially, prompting the LLM to explore and keep the INFO category in mind would have led to different results; we were therefore limited by keeping consistent prompting and not utilizing additional descriptive prompting.

Finally, disagreement with Reddit is not the same as moral error. AITA verdicts reflect the norms of a particular online community at a particular time, and reasonable moral pluralism can produce legitimate divergence. A model that declines to call both parties wrong may be expressing a defensible interpretive charity rather than failing to perceive shared fault. Our claim is not that the models are morally wrong where they diverge, but that the structure of their divergence reveals the dynamics of language modeling and alignment more than it reveals stable moral reasoning.

Future Work

Future work could build on this project by making the comparison between humans and models more direct. One useful extension would be to have individual human participants judge the same AITA posts under the same single-verdict format as the models, since this would make it clearer whether the patterns we found come from the models themselves or from the way Reddit comment threads produce consensus. Another important direction would be to test the same dilemmas under different relational, emotional, and personal-stakes framings to see whether the models hold onto the same verdicts when the moral facts stay the same but the conversational context changes. This would connect more directly to the question of moral stability that motivates the paper. Future work could also examine whether models with persistent memory

become more coherent across earlier and later judgments, or whether they simply appear more consistent because they are being reminded of their own outputs. Finally, because AITA reflects a mostly Western and English-speaking moral culture, it would be useful to test whether these same tendencies appear across different cultural contexts, especially in cases involving family obligation, privacy, loyalty, and responsibility.

Conclusion

Overall, both models are able to produce plausible moral judgments, but their divergences from human consensus are not the same. Across a dataset dominated by family-related disputes, humans, ChatGPT, and Claude often emphasized different parts of the moral situation. Human verdicts seemed more tied to fairness, reciprocity, and role-based obligations, while the LLMs more often leaned toward harm reduction and resolving conflict in a way that sounded socially appropriate or consequence-sensitive. ChatGPT tended to move toward shared or mutual fault, Claude tended to move toward more decisive single-party verdicts, and both models largely avoided acknowledging when there was not enough information to judge. This is especially clear in INFO cases, where the models almost always commit to a verdict rather than preserving uncertainty. These patterns matter because they show that model disagreements are not just different answers to the same questions, but may reflect different underlying moral emphases. At the same time, those emphases should not be mistaken for genuine moral commitments as the patterns we observe are better understood as products of training, alignment, and model-specific response tendencies than as evidence of moral reasoning in the human sense. Moral reasoning cannot be inferred from fluency or consistency alone so therefore, treating either as enough risks mistaking the simulation of moral judgment for the accountability, continuity, and self-correction that make moral reasoning meaningful in humans.

References

- “Am I the Asshole?” *Reddit*, <https://www.reddit.com/r/AmItheAsshole/>. Accessed 3 June 2026.
- ChatGPT. Response to "Moral Judgements in LLMs" OpenAI, 3 June 2026, <https://chatgpt.com/>
- Claude. Response to "Analyze Reddit AITA moral judgments." Anthropic, 3 June 2026, <https://claude.ai/>
- Haas, Julia. “A Theoretical Proposal for Modeling Artificial Moral Cognition.” *Minds & Machines*, vol. 30, 2020, pp. 219-246. *Springer Nature Link*, <https://doi.org/10.1007/s11023-020-09524-9>.
- Railton, Peter. “Ethical Learning, Natural and Artificial.” *Ethics of Artificial Intelligence*, vol. 17, 2020, pp. 45-78. *Oxford Academic*, <https://doi.org/10.1093/oso/9780190905033.003.0002>.
- Tomasello, Michael. *A Natural History of Human Morality*. Harvard University Press, 2018. Accessed 3 June 2026.
- Williams, Bernard. “Internal and External Reasons.” *Moral Luck: Philosophical Papers*, 1980, pp. 101-113. *Cambridge University Press*, <https://doi.org/10.1017/CBO9781139165860.009>.